

빅데이터 분석 및 응용

L01: MapReduce & DFS

Summer 2020

Kookmin University

In this lecture


- Motivation of MapReduce
- Distributed File System

MapReduce

- Large scale computing for data mining
- Challenges:
 - How to distribute computation?
 - Distributed/parallel programming is hard
- Map-reduce addresses all of the above
 - Google's computational/data manipulation model
 - Elegant way to work with big data

조작모델

Single Node Architecture




Motivation: Google Example

- 20+ billion web pages
- Average size of webpage = 20KB
- $20 \text{ billion} * 20\text{KB} = 400\text{TB}$
- Disk read bandwidth = 50MB/sec
- Time to read = 8 million seconds = 92+ days
- Even longer to do something useful with the data
- Today, a standard architecture for such problems is emerging:
 - Cluster of commodity Linux nodes
 - Commodity network (ethernet) to connect them

Cluster Architecture


1 Gbps between
any pair of nodes
in a rack



Each rack contains 16-64 nodes

In 2011 it was guesstimated that Google had 1M machines, <http://bit.ly/Shh0RO>

Cluster Architecture



Each rack contains 16-64 nodes

In 2011 it was guesstimated that Google had 1M machines, <http://bit.ly/Shh0RO>



Cluster Computing Challenges (1)

- **Node failures**
 - A single server may stay up for 3 years (1000 days)
 - If you have 1000 servers, expect to loose 1/day
 - People estimated Google has ~1M machines in 2011
 - 1000 machine fail every day!
- How to store data persistently and keep it available if nodes can fail?
- How to deal with node failures during a long-running computation?

Cluster Computing Challenges (2)

- **Network bottleneck**
 - Network bandwidth = 1 Gbps
 - Moving 10TB takes approximately 1 day
- **Distributed programming is hard!**
 - Need a simple model that hides most of the complexity

MapReduce

- MapReduce addresses the challenges of cluster computing
 - Store data redundantly on multiple nodes for persistence and availability
 - Move computation close to data to minimize data movement
 - Simple programming model to hide complexity of all this magic

Redundant Storage Infrastructure

- **Problem:**

- If nodes fail, how to store data persistently?

- **Answer:**

- **Distributed File System:**


- Provides global file namespace
- Google GFS; Hadoop HDFS;

- **Typical usage pattern**

- Huge files (100s of GB to TB)
- Data is rarely updated in place
- Reads and appends are common


Distributed File System

- Reliable distributed file system
- Data kept in “chunks” spread across machines



Distributed File System

- Reliable distributed file system
- Data kept in “chunks” spread across machines
- Each chunk **replicated** on different machines
 - Ensures persistence and availability



Chunk servers also serve as compute servers

Bring computation directly to the data!

Distributed File System

- **Chunk servers**

- File is split into contiguous chunks
- Typically each chunk is 16–64MB
- Each chunk replicated (usually 2x or 3x)
- Try to keep replicas in different racks

- **Master node**

- a.k.a. Name Node in Hadoop's HDFS
- Stores metadata about where files are stored
- Might be replicated

- **Client library for file access**

- Talks to master to find chunk servers
- Connects directly to chunk servers to access data

Questions?