

# 빅데이터 분석 및 응용

## L03: MapReduce Data Flow

Summer 2020

Kookmin University

# Map-Reduce: Environment

Map-Reduce environment takes care of:

- Partitioning the input data
- Scheduling the program's execution across a set of machines
- Performing the group by key step
- Handling machine failures
- Managing required inter-machine communication

# Map-Reduce: A diagram

## MAP:

Read input and produces a set of key-value pairs

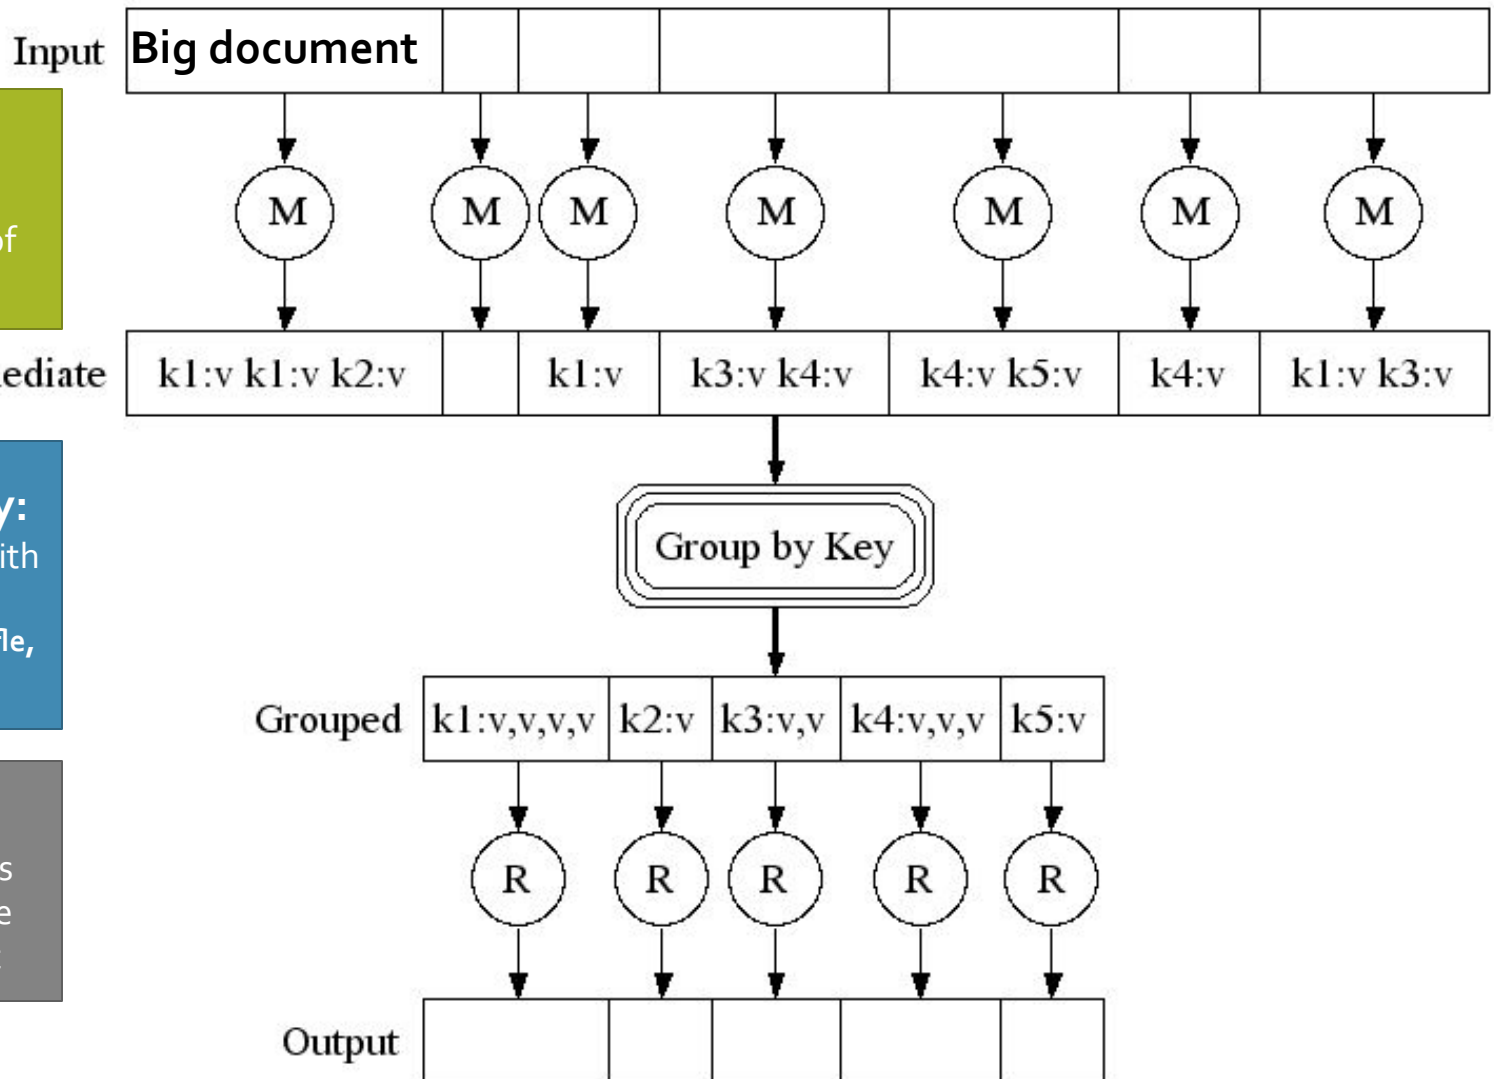
Intermediate

## Group by key:

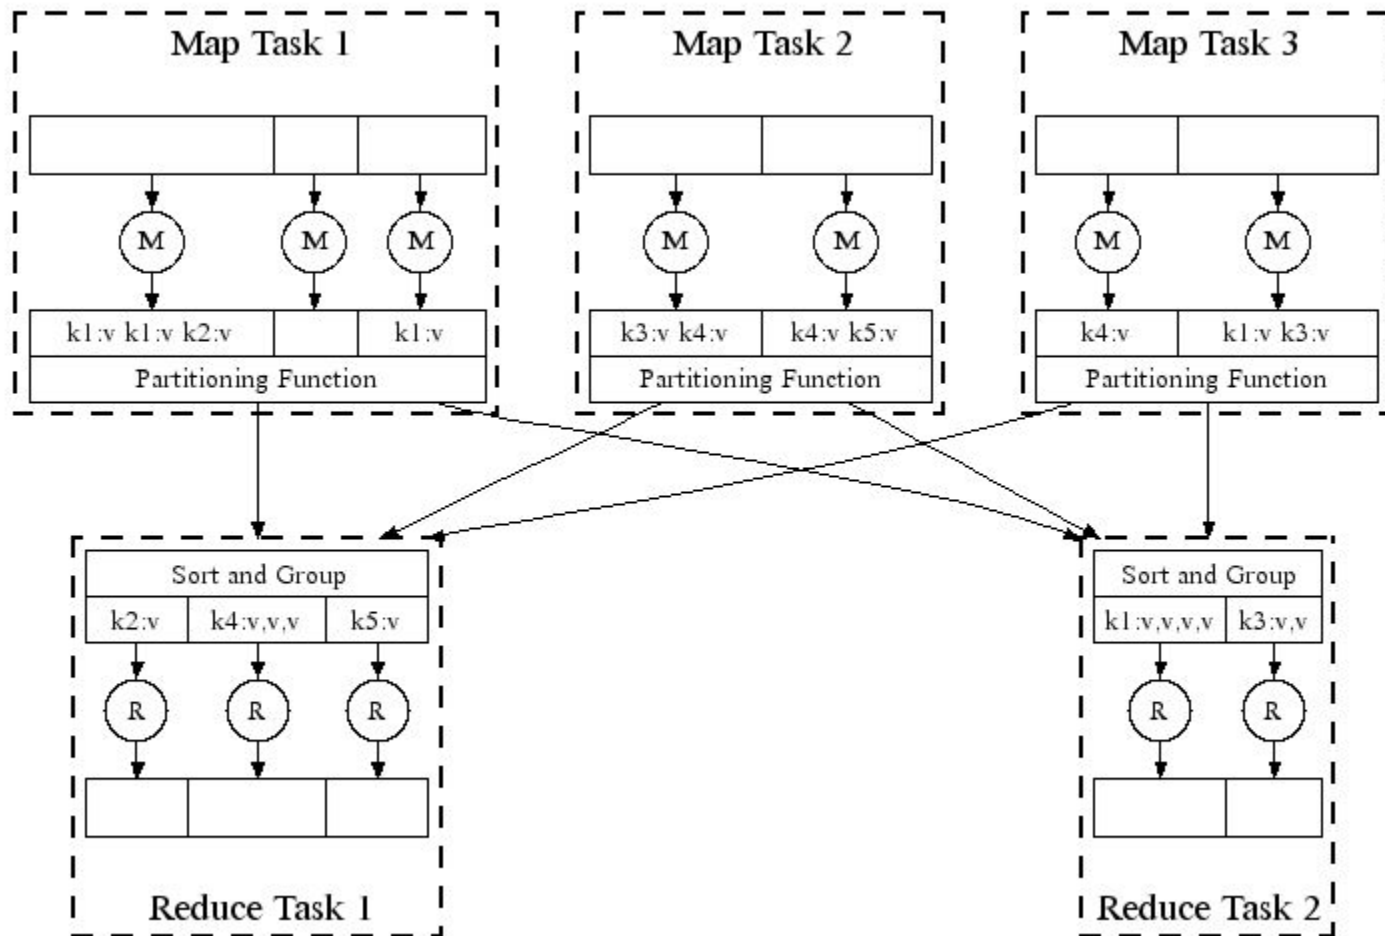
Collect all pairs with same key  
(Hash merge, Shuffle, Sort, Partition)

## Reduce:

Collect all values belonging to the key and output



# Map-Reduce: In Parallel



All phases are distributed with many tasks doing the work

# Data Flow

- Input and final output are stored on a distributed file system (DFS):
  - Scheduler tries to schedule map tasks “close” to physical storage location of input data
- Intermediate results are stored on local FS of Map and Reduce workers
- Output is often input to another MapReduce task

# Coordination: Master

- **Master node takes care of coordination:**
  - **Task status:** (idle, in-progress, completed)
  - **Idle tasks** get scheduled as workers become available
  - When a map task completes, it sends the master the location and sizes of its  $R$  intermediate files, one for each reducer
  - Master pushes this info to reducers
- Master pings workers periodically to detect failures

# Dealing with Failures

- **Map worker failure**

- Map tasks completed or in-progress at worker are reset to idle
- Reduce workers are notified when task is rescheduled on another worker

- **Reduce worker failure**

- Only in-progress tasks are reset to idle
- Reduce task is restarted

- **Master failure**

- MapReduce task is aborted and client is notified

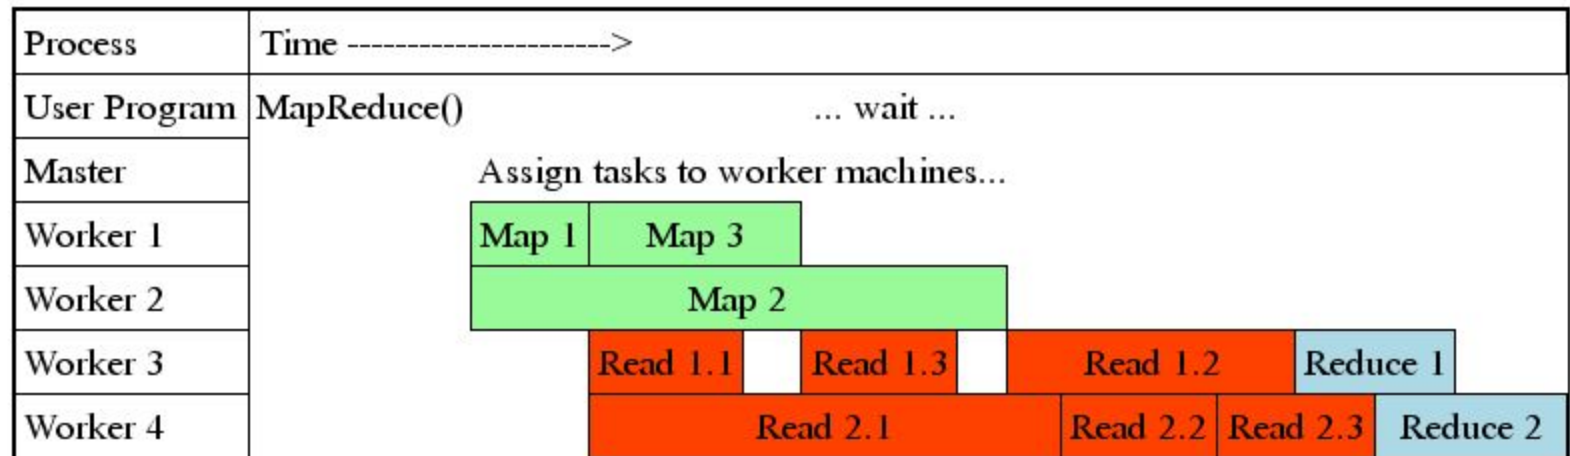
# How many Map and Reduce jobs?

- $M$  map tasks,  $R$  reduce tasks
- **Rule of a thumb:**
  - Make  $M$  much larger than the number of nodes in the cluster
  - One DFS chunk per map is common
  - Improves dynamic load balancing and speeds up recovery from worker failures
- **Usually  $R$  is smaller than  $M$** 
  - Because output is spread across  $R$  files



# Task Granularity & Pipelining

- **Fine granularity tasks:** map tasks  $\gg$  machines
  - Minimizes time for fault recovery
  - Can do pipeline shuffling with map execution
  - Better dynamic load balancing



# Refinement: Backup Tasks

- **Problem**

- Slow workers significantly lengthen the job completion time:
  - Other jobs on the machine
  - Bad disks
  - Weird things

- **Solution**

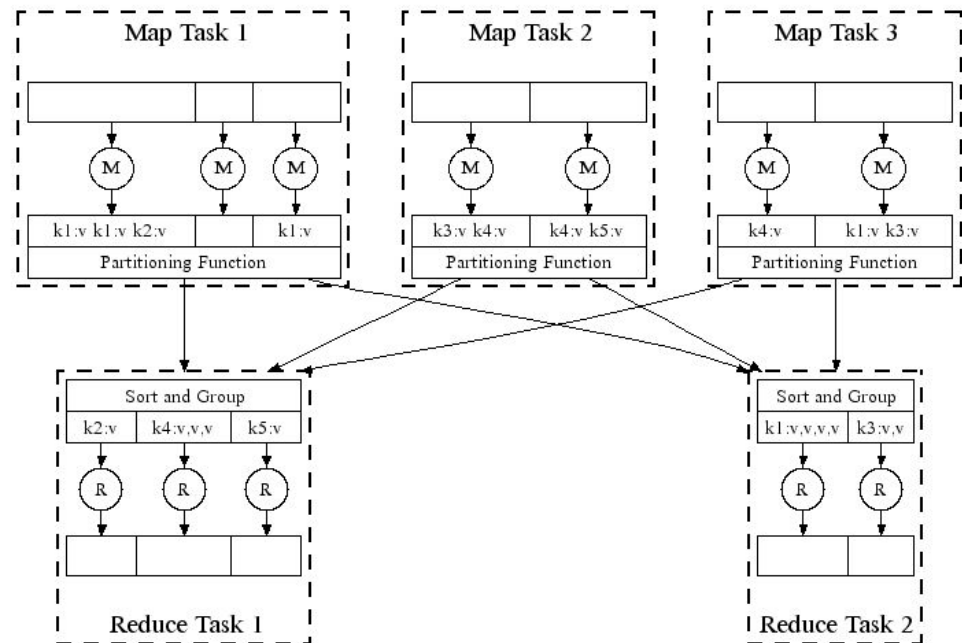
- Near end of phase, spawn backup copies of tasks
  - Whichever one finishes first “wins”

- **Effect**

- Dramatically shortens job completion time


# Refinement: Combiners

- Often a Map task will produce many pairs of the form  $(k, v_1), (k, v_2), \dots$  for the same key  $k$ 
  - E.g., popular words in the word count example
- Can save network time by **pre-aggregating values in the mapper:**
  - $\text{combine}(k, \text{list}(v_1)) \rightarrow v_2$
  - Combiner is usually same as the reduce function



# Refinement: Combiners

- **Back to our word counting example:**
  - Combiner combines the values of all keys of a single mapper (single machine):



# Refinement: Combiners

- Combiner trick works only if reduce function is commutative and associative

교환법칙

결합법칙

- Sum (o)**

- $a + b = b + a, (a + b) + c = a + (b + c)$

- Average ( $\Delta$ )**

- $\text{avg}(a,b) = \text{avg}(b,a), \text{avg}(\text{avg}(a,b), c) \neq \text{avg}(a, \text{avg}(b,c))$

결합 법칙 적용 불가

- Combination of sum and count

- Median (x)**

# Refinement: Partition Function

- Want to control how keys get partitioned
  - The set of keys that go to a single reduce worker
- System uses a default partition function:
  - $\text{hash}(\text{key}) \bmod R$
- Sometimes useful to override the hash function:
  - E.g.,  $\text{hash}(\text{hostname}(\text{URL})) \bmod R$  ensures URLs from a host end up in the same output file

# Implementations

- Google MapReduce
  - Uses Google File System (GFS) for stable storage
  - Not available outside Google
- Hadoop
  - Open-source implementation in Java
  - Uses HDFS for stable storage
  - Download: <http://hadoop.apache.org/>

# Questions?